

The Roots of Our Tax Debates

Matthew Weinzierl

AN AIR OF MYSTERY surrounds the recent debates over fiscal policy in Washington. How is it that compromise has proved so elusive? Considering that proposals from a diverse set of public and private fiscal commissions tend to overlap significantly, it seems obvious where the two sides should find middle ground: in the consensus that, to bring America's finances back into balance, spending must fall and revenues must rise.

The question of *how* to cut spending and raise revenues offers similar potential for agreement. On the spending side, non-defense discretionary spending has been repeatedly trimmed. If budgets are to be brought into balance through spending cuts, this leaves military spending (which has been reduced, though not as much as non-defense spending) and entitlement spending (which has been left largely untouched). On the revenue side, too, the list of plausible policy solutions is quite short, consisting at this point mostly of ways to broaden the tax base. Given that the options are thus limited, why is it so difficult to agree on some combination of these reforms?

The mystery deepens when we consider the stakes. Widely accepted projections show that America's long-term fiscal course is unsustainable. Deficits and debt are exceptionally high; income inequality raises concerns about the viability of the American Dream; and dissatisfaction with the tax code is nearly universal. Voters are aware of the problems and clearly want change. Legislators have even tried to tie their own hands through sequesters and other automatic budgeting mechanisms. Nevertheless, a grand bargain has failed to materialize.

MATTHEW WEINZIERL is an associate professor of business administration at Harvard Business School. This essay draws on his research with Benjamin B. Lockwood, N. Gregory Mankiw, and Danny Yagan.

But perhaps the absence of compromise is not so mysterious. Maybe the middle ground is not as obvious as it seems. After all, Americans are a large and diverse group of people. We differ in innumerable respects, prominent among them economic status, visions of the good life, and opinions about the role of government. If we all thought alike, policymaking would be simple. But since we do not—and since we often disagree about even the nature of the problems that public policy should aim to solve—reaching agreement takes time and debate. Because each of us often simultaneously accepts several different ideas about the *purpose* of public policy—to improve people's well-being, to reduce disparities of wealth, to advance growth and productivity, and so on—specific policy debates often involve prioritizing ends as much as choosing among means.

So while our fiscal debates appear superficially to be about means, might they in fact be the only way we can work through the problem of how to prioritize competing and at times even contradictory ends? If this is the case, what appears to be petty politicking over details may instead be a deeply consequential process—one that helps to define areas of agreement among widely differing sets of views about the proper purposes of public policy.

One particular aspect of the fiscal debate—the dispute over tax policy—offers the best example of how broader agreement is being prevented by differences over ends, not means, and how heated arguments about starkly different policy alternatives may actually be ways of seeking common ground.

For decades, these differences among us have been at the heart of how academic economists study tax policy. The insights this research is producing can help us to see past the surface of the familiar public debate and consider basic premises. A brief history of how economists have thought about the design of taxation—and of how the public understands tax policy—can thus help us see why our budget debates may be more constructive than they often seem.

MODELING OUR DIFFERENCES

One of the most important and highly regarded macroeconomic models was developed in the 1920s by a brilliant young mathematician named Frank Ramsey. The Ramsey model is an elegant and powerful simplification of macroeconomic reality. In it, a prototypical household decides

how hard to work and how much of its earnings to spend today or save and invest for the future, all with the goal of maximizing its well-being over time. In this model, a simplified government decides how best to tax that household in order to fund public expenditures. Today, economists use the Ramsey model to study fundamental questions like how economies grow.

For more than 80 years, the Ramsey model has also been a key tool for economists who study the design of tax policy. It has yielded some dramatic results, foremost among them the controversial and influential claim that capital income should not be taxed. Briefly, the idea is that taxing capital income reduces the amount of saving and investment and, therefore, the size of the capital stock that can be used in production. A smaller capital stock shrinks the economy so substantially under the model's assumptions that other taxes, such as those on labor income or consumption, are preferred.

But the source of the Ramsey model's appeal is also its greatest limitation: It assumes that all individuals are identical. In other words, it assumes that we can represent everyone in the economy with a single prototypical person (or household). Despite—and indeed because of—this extreme assumption, many economists argue that the Ramsey model remains a useful tool for answering some important questions. But when it comes to helping the country address complex political challenges like overhauling the tax system, the Ramsey model is insufficient.

Indeed, in the case of tax policy, the Ramsey model—rooted as it is in assumptions of uniformity—largely misses the point. Differences among individuals are at the heart of the tax-design problem. Not surprisingly, those differences are also at the heart of the other influential approach to understanding tax policy, the Mirrlees model. In 1971, economist James Mirrlees published an article in the *Review of Economic Studies* that changed how economists study taxes. At the core of the Mirrlees model is the recognition that individuals have different abilities to earn income. That is, for the same effort, some of us produce more income than others do. By carefully working out the implications of this simple insight, Mirrlees launched an intellectual revolution.

The Mirrlees model has yielded its own provocative claims, foremost among them that the ideal tax system would completely equalize after-tax incomes. To see where this claim comes from, we have to start by

understanding how this model defines the goal of tax policy. In the Mirrlees approach, the goal of tax policy is effectively to maximize overall well-being in society. Each individual has a level of well-being determined by how much leisure time and disposable income he has. Crucially, the well-being a person gets from an extra unit of leisure time or disposable income declines with how much of either he already has. In other words, a dollar is worth less to a rich person than to a poor one, such that redistributive taxes can raise overall well-being. In fact, the model suggests we should make such transfers until everyone gains the same well-being from an extra dollar of disposable income or hour of leisure time. Taking this idea to its logical conclusion, we ought to fully redistribute income from high earners to low earners. The ideal tax policy would be fully egalitarian.

Mirrlees's fame rests not on this provocative claim but on the way in which he showed why we cannot actually pursue this goal. The heart of the tax problem, as Mirrlees saw it, is that we cannot directly observe productivity differences among individuals. Instead, we see individuals' incomes, which are not the same as their levels of productivity. Incomes depend on productivity and effort, so when we tax people with high incomes more than we tax people with low incomes, we achieve some redistribution across productivity levels, but we also discourage effort. In more familiar terms, any success we have in dividing the pie more equally will come at the cost of having a smaller pie overall. By capturing this tradeoff between equality and efficiency, Mirrlees launched a new way of studying tax design. His model is the cornerstone of tax theory today.

But even the Mirrlees model ignores some important ways in which people differ. In fact, it ignores every difference other than the gaps in productivity. For example, it assumes that everyone places the same value on disposable income relative to leisure time. In reality, different attitudes toward what we call work-life balance among people with comparable natural abilities are commonplace. Similarly, the conventional Mirrlees approach assumes that everyone agrees that the goal of tax policy is to maximize overall well-being. But as our tax debates demonstrate, the proper goals of tax policy are still very much up for discussion in the United States.

It turns out that these other ways in which people differ have profound effects on how we ought to tax, and how we *do* tax.

TAXES AND THE GOOD LIFE

Part of the problem stems from a simple fact: High earners may be naturally talented, or hard-working, or, more often, some combination of both—and it is nearly impossible to isolate the true sources of their success. But does it matter whether someone earns a high income because he is talented or because he is hard-working? Put another way, does it matter whether someone earns a low income because he lacks talent or because he is unwilling to work hard? Most people think it does matter, and that a person should be held responsible for his willingness to work. After all, when we want to indicate our admiration for someone's success, we say, "He earned it." We may be wrong about that: It may be unjust to hold a person responsible for his work ethic, because the ability to persist through adversity and see a task to completion may be just as innate as any other skill, and therefore not really a function of character or will. But our belief in free will is powerful. Americans in particular celebrate the idea that hard work can bring success; when we see a high-income person, therefore, the possibility that he is a hard worker rather than a talented person makes us less willing to tax him for the benefit of others.

The claim that some people are willing to work hard means something quite specific to economists. It means that hard workers are more willing than others to give up leisure time in exchange for disposable income. In the context of the Mirrlees model, it means that people differ not only in their productivity but also in their tastes—in what they think makes for a good life.

If people differ in their tastes and not just in their abilities, then taxing high incomes and redistributing the revenue to others is not an ideal tool for maximizing overall well-being. Hard workers derive a lot of well-being from the last dollar they have earned, more perhaps than do people with lower incomes who have greater preferences for leisure. In that case, redistribution may even reduce overall well-being, especially if it has the side effect of shrinking the overall size of the pie.

The political philosopher Robert Nozick put this point especially vividly in his 1974 book, *Anarchy, State, and Utopia*: "Why should we treat the man whose happiness requires certain material goods or services differently from the man whose preferences and desires make such goods unnecessary for his happiness? Why should the man who prefers

seeing a movie (and who has to earn money for a ticket) be open to the required call to aid the needy, while the person who prefers looking at a sunset (and hence need earn no extra money) is not?" Nozick's concern was that an income tax — especially one that redistributes from those with high incomes to those with low incomes — assigns to two individuals two different levels of responsibility for providing for their fellow citizens, even if the only way in which the two differ is in how they want to spend their time.

This argument suggests that an awareness of different tastes and priorities among our fellow citizens will lead us to believe that, when a person has a high income, it is the result not merely of talent but of hard work. This, in turn, ought to make us less inclined to redistribute income through taxes.

The available evidence indeed confirms this expectation. The World Values Survey asks representative samples of individuals from a wide range of countries for their opinions on an array of questions pertaining to economics, politics, society, religion, and other topics. One of its questions asks respondents to say where they would place themselves on a work-leisure scale. This scale ranges from 1 to 5, where 1 corresponds to the statement "it's leisure that makes life worth living, not work" and 5 corresponds to the belief that "work is what makes life worth living, not leisure."

As the logic above would predict, in countries where answers vary more among respondents, standard measures of redistribution are statistically significantly lower. Sweden and the United States offer an illustrative comparison: The variance (a measure of the spread) of answers to the question above about the value of leisure versus work is about twice as large in the United States as it is in Sweden. Consistent with the logic that greater differences ought to lead to less redistribution, the ratio of government transfer payments (a common measure of redistribution) to total national economic output is more than twice as high in Sweden as it is in the United States. And this pattern holds more generally: Countries with less consensus about how to best use one's time, such as South Korea and Canada, have less redistributive tax policies than those with greater consensus, such as Finland and Germany.

Further evidence on this point comes from the states. The General Social Survey asks Americans how important it is to them to "have nice things." The survey uses a scale of 1 to 5; 1 means having nice things is

“one of the most important” priorities for a respondent, and 5 means it is “not at all important.” As is true among countries, the variation in answers to this question is small in some U.S. states and large in others. In those states with more variation, income taxes are less redistributive.

Thus both in theory and in fact, differences of opinion regarding what makes for a good life matter in designing tax policy. This lends further support to the notion that the sources of our fiscal debates may run deeper than we think—and that the differences across a vast and diverse population lie at the heart of our current disagreements.

THE HEIGHT TAX

Not all differences are created equal, however. As a business-school professor, I have lunches with small groups of students throughout the year. At some point during lunch, I usually ask: “Why do we tax?” I ask not to test the students—there is no “right” answer—but because I am interested in the diversity of their opinions. For some students, the answer is that we tax to maximize economic growth; others say the purpose of taxation is to maximize overall well-being; still others argue that we tax to provide equality of opportunity. As with differences in tastes, it turns out that understanding the differences of opinion regarding the purpose of taxation can help explain why we tax and what we tax.

The easiest way to illustrate the point is to propose a question that might, at first, seem absurd: Why don’t we levy a tax on height? To the modern tax economist, the idea of taxing height has a definite appeal. A height tax avoids the fundamental tradeoff between equality and efficiency—between slicing the pie evenly and keeping the pie large—that has forever challenged tax designers. Height is easy to verify and hard to change, at least relative to variables like income that existing tax systems target. And height is correlated with wages. The correlation is not strong, but it is there: On average, each additional inch of height for a white male in the United States corresponds to an increase in income of about 2%. By taxing height, we would transfer resources from, on average, higher-income to lower-income people, with minimal negative effects on effort (because the tall cannot avoid the height tax by working less).

The source of the height premium in wages has been extensively studied, and the results make taxing height even more appealing. While the full explanation is complicated, the most compelling evidence suggests

that the height premium may be traced to factors such as prenatal health (which is positively related to both adult height and cognitive ability) that are innate to each individual, rather than acquired or earned. In other words, taxing height would help us to reduce the fundamental barrier to optimal taxation that Mirrlees identified — namely, that we cannot observe innate productivity.

But of course very few of us would actually support taxing height, even if these facts were explained to us. When I ask my students why we don't tax height, they give answers such as "it seems arbitrary" or "we simply don't allow that sort of discrimination." Those answers feel right, but they don't help us very much. The benefit of taxing height is clear: improved equality of outcomes without many efficiency costs (if any at all). To explain why we intuitively reject height taxes, then, we must make the downsides equally clear.

The stakes involved in explaining why we reject height taxes are large. After all, the recommendation to tax height derives from the central premises of the Mirrlees model that tax theorists use to understand and design policy. So if we find height taxes unpalatable, in what other ways is that same model yielding policy recommendations that are at odds with our society's preferences? Is that model a reliable guide to policy design, or does it need to be replaced?

In particular, our rejection of a height tax should prompt us to question why we endorse substantial income redistribution through the tax code. The height tax largely addresses the classic problem with redistribution — namely, the dulling of incentives. If we reject a redistributive tax that has almost no incentive costs, how can we be in favor of redistribution that does generate such costs? And, if we are in favor of such redistribution, is it for reasons other than the maximization of well-being, which most tax theorists assume is our goal?

Complicating matters is the fact that, in the United States, we do base taxes in part on some personal characteristics. The example most readily comparable to height is sight. The tax credit for those with limited sight implicitly raises taxes on other taxpayers (albeit by very little, given the credit's small size and scope). Most of us find this tax credit entirely reasonable: When my students are asked why a sight tax is acceptable, they note that blindness seems a much better predictor of having difficulty earning a high income. (Official data on the incomes of those who claim the blindness tax credit are certainly consistent with

this impression.) Therefore, any explanation for why we avoid taxing height must be consistent with our embrace of taxing sight.

One answer to the height-tax puzzle has its roots in tax theory from the 19th century. John Stuart Mill, an economist and moral philosopher, is famous in part as a founder of utilitarianism, the philosophy that assumes the goal of society is to maximize overall well-being. This is the philosophy upon which the Mirrlees model is based.

Ironically, however, Mill did not embrace utilitarianism as the proper guide to tax policy. Instead, he thought that taxes ought to be assigned so that all individuals bear the same sacrifice—or loss of well-being—by paying them. This principle of equal sacrifice was highly influential for many decades, though its influence waned with the rise of the Mirrlees approach.

If taxes followed the principle of equal sacrifice, we would not tax height. Take two people who earn the same income, distinguished by the fact that one is short and the other is tall. It seems reasonable to assume that, at least on average, two such people put forth the same effort to earn their incomes, which means they have the same level of productivity. If we were to tax height, the taller person would pay more than the short person. In other words, we would ask for greater sacrifice from the tall person than from the short person.

But if we followed the principle of equal sacrifice rigorously, we would also avoid taxing sight. Just as taxing the tall more than the short would impose unequal sacrifice, so too would taxing those who can see more than those who are blind. Because we reject height taxes while accepting sight taxes, the explanation for this double standard can't be quite so simple (or radical) as choosing equal sacrifice as an alternative principle of tax design.

As with the earlier discussion of taste differences, one way to resolve this puzzle involves recognizing that people differ in the principles they believe ought to guide tax policy. Suppose that some people accept wholeheartedly the utilitarian idea that the goal of tax policy is to maximize overall well-being. They would endorse both the height tax and the sight tax; they would not mind the unequal sacrifice these taxes would involve, because they place a high value on the redistribution these taxes would achieve. Suppose that everyone else, however, is of the opposite opinion. These people accept wholeheartedly the idea that the goal of tax policy is to impose equal sacrifice on everyone. They would reject both the height

tax and the sight tax—placing a high value on protecting equal sacrifice and not valuing the redistribution these taxes achieve.

Now suppose people have to negotiate over what policies they will choose as a society. If either of the aforementioned groups dominates, we should see both a height tax and a sight tax, or neither. But can these people find a middle ground?

Yes, and the middle ground is to reject a height tax but accept a sight tax—just what we do in the United States today. To understand why, note that the relationship between height and income is weaker than the relationship between sight and income. This minimizes the benefits and magnifies the costs of the height tax: Transfers from the tall to the short will take from many low earners and give to many high earners, because height is such a weak predictor of income. The “unintended” transfers under a height tax reduce its value in terms of redistribution and increase its cost in terms of unequal sacrifice. By contrast, transfers from people with sight to the blind will largely benefit individuals with low incomes, so that redistribution will be achieved with only minor violations of the equal-sacrifice principle.

But the thought experiment outlined above is not quite relevant to the question at hand. I have asked thousands of people which of these principles—maximized well-being or equal sacrifice—they believe should govern tax policy. With striking regularity, only about one-third of those surveyed endorse either principle wholeheartedly. Two-thirds prefer a mixture of the two.

In other words, a large majority of people end up negotiating with *themselves* over which principle of tax policy they endorse. This makes sense: Both principles have appeal, and neither is fully convincing. Most of us are drawn to the idea that because some people are privileged in ways that many others are not, the lucky should therefore share their good fortune. We are also often drawn to the idea that society is a voluntary arrangement in which each member may be asked to contribute equally to the funding of common activities. In the end, then, many of us prefer a policy that gives some weight to each set of priorities. Evidence suggests that, when judging tax policy, most of us weigh the benefits in terms of increased social well-being and the costs in terms of unequal sacrifice.

Appreciating this internal ambivalence over the goals of tax policy is a powerful tool—not only for understanding the policies that currently

govern our economy and society, but also for understanding the passionate debates in Washington that, over the past several years, have prevented much-needed policies from being enacted.

FINDING A BALANCE

Countless commentators have lamented Congress's seeming inability to find a middle ground in the debates over taxation. The ferocity of the battles over the debt limit, the fiscal cliff, the top marginal tax rate, and entitlement reforms are seen as discouraging signs that the great deliberative bodies of the legislative branch are broken.

This conventional analysis misses a more encouraging explanation suggested by the discussion above: that Americans are wrestling with the profound question of why we tax. If we interpreted the tax debate as helping voters decide what they believe ought to guide tax policy, then we would see the potential for these bitter battles to yield, in the end, a more robust compromise. We would also understand that those most convinced of their perspectives—say, party officials and activists—see the current moment as an opportunity to convince others that the right balance of priorities puts more weight on theirs.

In other words, the battles we have seen can be interpreted not as evidence that politicians won't meet in the middle, but rather as an indication that politicians are trying to affect where the middle is—where we strike the balance between the benefits of social well-being and the costs of unequal sacrifice. They are seeking to push each individual, and thereby the population as a whole, toward their own positions. They understand that they will fully convince almost no one, but that's not the point. If most of our opinions represent some balance among multiple principles, parties and politicians can make progress simply by altering the weight we put on each one. Arguably, achieving that goal would yield dramatically greater long-term benefits for any political party than would appealing to moderate voters by accepting a compromise.

The rhetoric that Barack Obama and Mitt Romney used in the 2012 presidential election suggests that just such a battle was being waged. Each candidate forcefully argued that this was an election about what we wanted the role of the federal government—and the guiding principles of our economy—to be. As Peter Baker wrote in the *New York Times* after the first presidential debate, when "President Obama and Mitt Romney faced off for the first time, their largely zinger-free styles

may have disguised a fierce clash of views not only over taxes, spending and health care, but over the very role of government in American society in a time of wrenching problems.”

Though the 2012 election feels like a distant memory, we have not left those fierce clashes behind. The recent debate over the fiscal year 2014 budget illustrates the point. While both the Republican-led House and the Democratic-led Senate passed budget proposals earlier this spring, they are so far apart that the conventional next step—formally reconciling the proposals in a conference committee—is widely seen as unlikely to succeed. The reason for the impasse is clear: Senate majority leader Harry Reid described the House plan as “an extreme budget that is anything but balanced. . . . [It] reflects the same skewed priorities the Republican party has championed for years—the same skewed priorities Americans rejected in November.” Aside from the adjectives, House Budget Committee chairman Paul Ryan would not necessarily disagree. When announcing the House proposal, he said: “The election didn’t go our way. . . . That means we surrender our principles? That means we stop believing in what we believe in?”

The two sides of this debate would seem to be insisting on absolute victory. They repeatedly rehearse the arguments for less spending or more taxes to solve our fiscal imbalances, the cases for flatter tax rates to encourage growth and individual effort or greater redistribution for the sake of fairness and general well-being, and so on. The key to understanding these arguments, however, is to acknowledge that each party is trying to influence the public’s balance of priorities, not to eliminate that balance in favor of an absolute preference.

It may look as if the debates over American fiscal policy are going nowhere. But if we recognize the central role that differences among us play in how we pursue compromise in the design of taxes, we can appreciate how these arguments make progress possible. Debates over such fundamental questions *ought* to be passionate. They encourage each of us to examine whether we understand other people’s positions. And because most of us find a number of positions tenable, these discussions encourage us to more thoroughly examine our own positions, too. In the end, we will meet in the middle; first, however, we must decide where that middle is.

Copyright of National Affairs is the property of National Affairs, Inc. and its content may not be copied or emailed to multiple sites or posted to a listserv without the copyright holder's express written permission. However, users may print, download, or email articles for individual use.